

LLM Policy

Ethan Levien

The purpose of this document is to state the course policy, its rationale, and provide some guidance on the use of LLMs¹

Course Policy

You are free to use LLMs at will on all take-home assignments, subject to the following:

- All LLM output must be reviewed and edited. If I feel I am reading completely LLM-generated answers with long, padded sentences and LLM tropes, I reserve the right to deduct points, even if they are correct.
- When you use an LLM for an assignment, you must include an *LLM disclosure statement* with the following information:
 - What models were used and how (e.g., Claude Code vs. a chat app).
 - A 1-2 paragraph reflection summarizing how the model was used: List the main contributions (e.g., pieces of code, making a specific plot) and how you engaged with those contributions to ensure the assignment was still beneficial.

Finally, unless otherwise indicated, it should be assumed that an assignment can be completed without an LLM.

You are not permitted to use LLMs or any electronics during in-class exams or during lecture.

Rationale

The point that, unless otherwise indicated, an assignment can be completed without an LLM is emphasized in the course policy because I feel it is the most important part. I am well aware that most of you are using LLMs heavily, as am I. It seems as if this is inevitable in the current moment. We often find ourselves in situations where it is difficult to justify doing a task “by hand” when it can be done faster and, in many instances, more accurately with the assistance of an LLM. Moreover, the idea that being able to use LLMs effectively is key to your professional success is widespread (and might be true). In mathematics, we have seen many examples where top researchers are partnering with LLM companies to obtain high-profile results [2]. These give us the impression that “AI is the future” and every endeavor should maximally leverage these tools.

For me personally, LLMs have been extremely beneficial and accelerated certain aspects of my research. The experience of being able to have an intelligent conversation about an obscure

¹I use the term LLM rather than AI in the course policy because when we refer to AI, we are usually referring to a set of Large Language Models (LLMs) developed by big tech companies like OpenAI, Google, and Anthropic. I consider AI to be a broader set of ideas – see [1].

and technical research topic with a computer is quite frankly unbelievable. At the same time, they can be incredibly destructive to our ability to learn and think independently when used improperly. This is not to mention the immense social and environmental harm. It is unclear that the claimed economic benefits will outweigh these harms. Many people have written about these concerns, and I will not go deeply into them here:

- Michael Harris, in his newsletter *Silicon Reckoner* [3]
- Emily Bender, in *The AI Con* [4]
- Carl T. Bergstrom and Jevin D. West, *Modern-Day Oracles or Bullshit Machines?* [5]
- Research on AI in education [6]
- A very long list of AI critical articles [7]

I recommend you read at least some portion of something on that list.

The course policy tries to strike a balance between a highly skeptical view and being practical: Yes, LLMs are important and using them is a skill you will benefit from developing, but you should not be required to use them at this level and if you choose to use them, it should be done with an awareness of the risks.

Recommendations

Here are some general practical recommendations for those that decide to use LLMs in this course:

- Find a way to document exactly how AI has contributed to a project. I like to use git² on a research project and have separate commits for things I have done vs. Claude. I can then go through the history and see exactly which things were injected by the LLM (and therefore need to be reviewed). Git might be overkill for course work, but you should consider saving different versions of a document (e.g., before and after AI proofreading).
- Set a time constraint for thinking about a problem without an LLM. It may feel like a way to get unstuck: you’ve been thinking about a problem for one minute, then you throw it at an LLM. Were you really stuck? Did you try everything? Give yourself at least 10 minutes for each drill. The experience of being stuck, sometimes for long periods of time, is a central part of doing science. It is an experience you need to practice and become used to.
- Don’t judge your productivity by intermediate steps. LLMs are very effective at getting you through certain stages of a project/assignment, but are they really helping you be more productive in the long term? It may seem like you are saving time by using LLMs for “trivial” things so that you can “get to the interesting part.” An example might be working through some algebra to arrive at the more conceptually interesting part of a problem. But ask yourself, could doing that seemingly trivial task actually be helping your brain make connections that shape how you approach the more conceptual aspects?

²<https://git-scm.com/>

- Notice if there are tasks you have not done for some time without an LLM. For example, you may get used to using an LLM to turn bullet points into finished prose. When was the last time you actually wrote something down from start to finish with complete sentences and correct grammar? Make sure to practice the skills you want to maintain.

References

- [1] Emily M. Bender. Artificial intelligence. In Rayvon Fouché, editor, *Oxford Research Encyclopedia of Science, Technology, and Society*. Oxford Academic, New York, NY, June 2026. Online edn., 25 June 2026–, <https://doi.org/10.1093/9780197852712.003.0067>.
- [2] Noga Alon, Thomas F. Bloom, W. T. Gowers, Daniel Litt, Will Sawin, Arul Shankar, Jacob Tsimmerman, Victor Wang, and Melanie Matchett Wood. Remarks on the disproof of the unit distance conjecture, 2026. A human-verified account of an OpenAI-generated disproof, with reflections by the authors, <https://arxiv.org/abs/2605.20695>.
- [3] Michael Harris. Silicon reckoner. Substack newsletter, <https://siliconreckoner.substack.com/>.
- [4] Emily M. Bender and Alex Hanna. *The AI Con: How to Fight Big Tech’s Hype and Create the Future We Want*. Harper, New York, NY, 2025.
- [5] Carl T. Bergstrom and Jevin D. West. Modern-day oracles or bullshit machines? how to thrive in a chatgpt world. <https://thebullshitmachines.com/>.
- [6] David Strömberg, Victor Lei, and Yanhui Wu. The generative ai learning penalty: Evidence from chinese secondary education. Technical report, SSRN, June 2026. <https://ssrn.com/abstract=6868618>.
- [7] Mirabelle Jones. Stochastic parrots reading / viewing list. Google Docs, compiled for Stochastic Parrots Day, DAIR, <https://docs.google.com/document/d/1bG0yIdawiUvwh7m0AnXV5W6JHkK9xwXemuVjSU5tbhQ/>, 2023.