

Lecture 10 (10/7)

Confidence intervals, Least squares

Summary of material since midterm 1:

- Unit 2.4: Correlation (ρ) as a dimensionless way to measure relatedness between variables. Coefficient of determination ρ^2 is the fraction of variance explained by X .

Summary of material since midterm 1:

- Unit 2.4: Correlation (ρ) as a dimensionless way to measure relatedness between variables. Coefficient of determination ρ^2 is the fraction of variance explained by X .
- Unit 3.2: Estimators are approximations of “ground truth” from data vary from dataset to dataset. This variation is modeled with the sample distribution.

Summary of material since midterm 1:

- Unit 2.4: Correlation (ρ) as a dimensionless way to measure relatedness between variables. Coefficient of determination ρ^2 is the fraction of variance explained by X .
- Unit 3.2: Estimators are approximations of “ground truth” from data vary from dataset to dataset. This variation is modeled with the sample distribution.
- Unit 3.2: Error in estimators is decomposed into Bias and variance:

$$\text{MSE}_{\hat{\theta}} = \text{se}(\hat{\theta})^2 + \text{Bias}^2.$$

For a consistent estimator both shrink with the amount of data we collect.

Summary of material since midterm 1:

- Unit 2.4: Correlation (ρ) as a dimensionless way to measure relatedness between variables. Coefficient of determination ρ^2 is the fraction of variance explained by X .
- Unit 3.2: Estimators are approximations of “ground truth” from data vary from dataset to dataset. This variation is modeled with the sample distribution.
- Unit 3.2: Error in estimators is decomposed into Bias and variance:

$$\text{MSE}_{\hat{\theta}} = \text{se}(\hat{\theta})^2 + \text{Bias}^2.$$

For a consistent estimator both shrink with the amount of data we collect.
Recommended exercises (in addition to all drills): 2.5, 3.6, 3.7.

Summary of material since midterm 1:

- Unit 2.4: Correlation (ρ) as a dimensionless way to measure relatedness between variables. Coefficient of determination ρ^2 is the fraction of variance explained by X .
- Unit 3.2: Estimators are approximations of “ground truth” from data vary from dataset to dataset. This variation is modeled with the sample distribution.
- Unit 3.2: Error in estimators is decomposed into Bias and variance:

$$\text{MSE}_{\hat{\theta}} = \text{se}(\hat{\theta})^2 + \text{Bias}^2.$$

For a consistent estimator both shrink with the amount of data we collect.
Recommended exercises (in addition to all drills): 2.5, 3.6, 3.7.

Today:

- Unit 3.3/3.6: Confidence intervals, using the estimated sample distribution, including Hypothesis testing.
- Unit 3.4: Least squares estimators of $\beta_0, \beta_1, \sigma_\epsilon^2$ and R^2 , and the sample distribution of $\hat{\beta}_1$.

True vs. estimated sample distribution

For $Y_i \sim \text{Normal}(\mu, \sigma^2)$, the *true* sample distribution of $\hat{\mu} = \bar{Y}$ is

$$\hat{\mu} \sim \text{Normal}\left(\mu, \frac{\sigma^2}{N}\right).$$

True vs. estimated sample distribution

For $Y_i \sim \text{Normal}(\mu, \sigma^2)$, the *true* sample distribution of $\hat{\mu} = \bar{Y}$ is

$$\hat{\mu} \sim \text{Normal}\left(\mu, \frac{\sigma^2}{N}\right).$$

But we don't know μ or σ ! In practice we plug in our estimates to get the estimated sample distribution

$$\hat{\mu} \sim \text{Normal}\left(\hat{\mu}, \frac{\hat{\sigma}^2}{N}\right) \quad (\text{red: numbers computed from our data})$$


True vs. estimated sample distribution

For $Y_i \sim \text{Normal}(\mu, \sigma^2)$, the *true* sample distribution of $\hat{\mu} = \bar{Y}$ is

$$\hat{\mu} \sim \text{Normal}\left(\mu, \frac{\sigma^2}{N}\right).$$

But we don't know μ or σ ! In practice we plug in our estimates to get the estimated sample distribution

$$\hat{\mu} \sim \text{Normal}\left(\hat{\mu}, \frac{\hat{\sigma}^2}{N}\right) \quad (\text{red: numbers computed from our data})$$



Points: data ($N = 10$). Black: true. Red dashed: estimated, centered at $\hat{\mu}$.


True vs. estimated sample distribution

For $Y_i \sim \text{Normal}(\mu, \sigma^2)$, the *true* sample distribution of $\hat{\mu} = \bar{Y}$ is

$$\hat{\mu} \sim \text{Normal}\left(\mu, \frac{\sigma^2}{N}\right).$$

But we don't know μ or σ ! In practice we plug in our estimates to get the estimated sample distribution

$$\hat{\mu} \sim \text{Normal}\left(\hat{\mu}, \frac{\hat{\sigma}^2}{N}\right) \quad (\text{red: numbers computed from our data})$$



Points: data ($N = 10$). Black: true. Red dashed: estimated, centered at $\hat{\mu}$.

We need to estimate the parameter in the sample distribution, which means the entire distribution is itself a random variable. Still, its width tells us how far μ might plausibly be from $\hat{\mu}$. This idea is formalized with confidence intervals.

Confidence intervals

The 95% confidence interval:

$$\text{CI} = \left[\hat{\theta} - 1.96 \text{se}(\hat{\theta}), \hat{\theta} + 1.96 \text{se}(\hat{\theta}) \right].$$

Confidence intervals

The 95% confidence interval:

$$\text{CI} = \left[\hat{\theta} - 1.96 \text{se}(\hat{\theta}), \hat{\theta} + 1.96 \text{se}(\hat{\theta}) \right].$$


For an approximately Normal sample distribution, 95% of values fall within 1.96 standard errors of the center.

Confidence intervals

The 95% confidence interval:

$$\text{CI} = \left[\hat{\theta} - 1.96 \text{se}(\hat{\theta}), \hat{\theta} + 1.96 \text{se}(\hat{\theta}) \right].$$

For an approximately Normal sample distribution, 95% of values fall within 1.96 standard errors of the center.



Each bar is the CI from one replicate, with a dot at $\hat{\theta}$. About 95% overlap with θ .

Subtleties of interpretation (optional)

If you want to get pedantic about CIs, there is a distinction which is more about the language we use to talk about them than how they are used practically:

- **Correct:** if we generate many replicates of the data, the true θ falls in the resulting CI for 95% of them; that is,

$$P(\theta \in CI) = .95$$

where CI is treated as a random variable.

Subtleties of interpretation (optional)

If you want to get pedantic about CIs, there is a distinction which is more about the language we use to talk about them than how they are used practically:

- **Correct:** if we generate many replicates of the data, the true θ falls in the resulting CI for 95% of them; that is,

$$P(\theta \in \text{CI}) = .95$$

where CI is treated as a random variable.


- **Not correct:** there's a 95% chance θ lies in *this particular* interval; that is, say we compute $\text{CI} = [0.2, 1.4]$.

$$P(\theta \in \text{CI}) = 0 \quad \text{OR} \quad P(\theta \in \text{CI}) = 1$$

Example

Midterm 1 was graded by two graders, E and G. Let Y be the score (out of 25) and $X = 1$ if E graded the exam and model the scores as

$$Y \mid X \sim \text{Normal}(\beta_0 + \beta_1 X, \sigma_\epsilon^2).$$



- Compute $\overline{Y|X=0}$, $\overline{Y|X=1}$, $\hat{\beta}_1$ and $\hat{\sigma}_\epsilon^2$.
- What are the sample distributions $\hat{\beta}_1$?

Sample distribution of a difference of means

An “eyeball” calculation

$$\overline{Y|X=0} \approx 17, \quad \overline{Y|X=1} \approx 17.5, \quad \hat{\sigma}_\epsilon \approx 5$$

hence

$$\overline{Y|X=1} \sim \text{Normal}(17.5, 25/23), \quad \overline{Y|X=0} \sim \text{Normal}(17, 25/45),$$

and therefore using $25/23 + 25/45 \approx 1 + 1/2 = 3/2$,

$$\hat{\beta}_1 = \overline{Y|X=1} - \overline{Y|X=0} \sim \text{Normal}(1/2, 3/2)$$

Sample distribution of a difference of means

An “eyeball” calculation

$$\overline{Y|X=0} \approx 17, \quad \overline{Y|X=1} \approx 17.5, \quad \hat{\sigma}_\epsilon \approx 5$$

hence

$$\overline{Y|X=1} \sim \text{Normal}(17.5, 25/23), \quad \overline{Y|X=0} \sim \text{Normal}(17, 25/45),$$

and therefore using $25/23 + 25/45 \approx 1 + 1/2 = 3/2$,

$$\hat{\beta}_1 = \overline{Y|X=1} - \overline{Y|X=0} \sim \text{Normal}(1/2, 3/2)$$

In Python (with `x`, `y` numpy arrays):

```
y0 = y[x == 0]
y1 = y[x == 1]
b1_hat = np.mean(y1) - np.mean(y0)
var_b1 = np.std(y1)**2/len(y1) + np.std(y0)**2/len(y0)
```

gives us

$$\hat{\beta}_1 = \overline{Y|X=1} - \overline{Y|X=0} \sim \text{Normal}(0.68, 1.84).$$

Hypothesis test

Example

Using “eyeball” numbers

- Find the 95%-CI for β_1 .
- Hypothesis test: if the null hypothesis $H_0 : \beta_1 = 0$ (same grading) were true, the test statistic $\hat{T}^* = \hat{\beta}_1^* / \text{se}(\hat{\beta}_1) \sim \text{Normal}(0, 1)$. To answer the question: Can we reject this hypothesis, we compute p-value $P(|\hat{T}^*| > |\hat{T}|)$.

Hypothesis test

- We already found $\text{se}(\hat{\beta}_1) \approx \sqrt{3/2} \approx 1.2$, so (using $1.96 \approx 2$)
 $\text{CI} \approx [1/2 - 2(1.2), 1/2 + 2(1.2)] \approx [-2, 3]$, which contains 0.

Hypothesis test


- We already found $\text{se}(\hat{\beta}_1) \approx \sqrt{3/2} \approx 1.2$, so (using $1.96 \approx 2$)
 $\text{CI} \approx [1/2 - 2(1.2), 1/2 + 2(1.2)] \approx [-2, 3]$, which contains 0.
- The test statistic is $\hat{T} = \hat{\beta}_1 / \text{se}(\hat{\beta}_1) \approx (1/2) / 1.2 \approx 0.4$: our difference is only 0.4 standard errors from 0.

Hypothesis test

- We already found $\text{se}(\hat{\beta}_1) \approx \sqrt{3/2} \approx 1.2$, so (using $1.96 \approx 2$)
 $\text{CI} \approx [1/2 - 2(1.2), 1/2 + 2(1.2)] \approx [-2, 3]$, which contains 0.
- The test statistic is $\hat{T} = \hat{\beta}_1/\text{se}(\hat{\beta}_1) \approx (1/2)/1.2 \approx 0.4$: our difference is only 0.4 standard errors from 0.
- Near the center the Normal density is about 0.4, so $P(|\hat{T}^*| < 0.4) \approx \approx 0.3$ (from bell curve), hence $p_v = P(|\hat{T}^*| > 0.4) \approx 0.7$. A difference like ours is typical if the graders grade the same, so we *cannot* reject H_0 .

Hypothesis test

- We already found $\text{se}(\hat{\beta}_1) \approx \sqrt{3/2} \approx 1.2$, so (using $1.96 \approx 2$) $\text{CI} \approx [1/2 - 2(1.2), 1/2 + 2(1.2)] \approx [-2, 3]$, which contains 0.
- The test statistic is $\hat{T} = \hat{\beta}_1/\text{se}(\hat{\beta}_1) \approx (1/2)/1.2 \approx 0.4$: our difference is only 0.4 standard errors from 0.
- Near the center the Normal density is about 0.4, so $P(|\hat{T}^*| < 0.4) \approx \approx 0.3$ (from bell curve), hence $p_v = P(|\hat{T}^*| > 0.4) \approx 0.7$. A difference like ours is typical if the graders grade the same, so we *cannot* reject H_0 .



Estimators of regression parameters

Data $(X_1, Y_1), \dots, (X_N, Y_N)$ from $Y | X \sim \text{Normal}(\beta_0 + \beta_1 X, \sigma_\epsilon^2)$. How do we estimate β_1 ?

Estimators of regression parameters

Data $(X_1, Y_1), \dots, (X_N, Y_N)$ from $Y | X \sim \text{Normal}(\beta_0 + \beta_1 X, \sigma_\epsilon^2)$. How do we estimate β_1 ?

We know

$$\beta_1 = \frac{\text{cov}(X, Y)}{\sigma_X^2}.$$

Estimators of regression parameters

Data $(X_1, Y_1), \dots, (X_N, Y_N)$ from $Y | X \sim \text{Normal}(\beta_0 + \beta_1 X, \sigma_\epsilon^2)$. How do we estimate β_1 ?

We know

$$\beta_1 = \frac{\text{cov}(X, Y)}{\sigma_X^2}.$$

To estimate $\text{cov}(X, Y)$ from data, it helps to write it in a different way. Start from

$$E[(X - E[X])(Y - E[Y])]$$

Estimators of regression parameters

Data $(X_1, Y_1), \dots, (X_N, Y_N)$ from $Y \mid X \sim \text{Normal}(\beta_0 + \beta_1 X, \sigma_\epsilon^2)$. How do we estimate β_1 ?

We know

$$\beta_1 = \frac{\text{cov}(X, Y)}{\sigma_X^2}.$$

To estimate $\text{cov}(X, Y)$ from data, it helps to write it in a different way. Start from

$$\begin{aligned} E[(X - E[X])(Y - E[Y])] \\ = E[XY - E[Y]X - E[X]Y + E[X]E[Y]] \quad (\text{expand}) \end{aligned}$$

Estimators of regression parameters

Data $(X_1, Y_1), \dots, (X_N, Y_N)$ from $Y | X \sim \text{Normal}(\beta_0 + \beta_1 X, \sigma_\epsilon^2)$. How do we estimate β_1 ?

We know

$$\beta_1 = \frac{\text{cov}(X, Y)}{\sigma_X^2}.$$

To estimate $\text{cov}(X, Y)$ from data, it helps to write it in a different way. Start from

$$\begin{aligned} & E[(X - E[X])(Y - E[Y])] \\ &= E[XY - E[Y]X - E[X]Y + E[X]E[Y]] \quad (\text{expand}) \\ &= E[XY] - E[Y]E[X] - E[X]E[Y] + E[X]E[Y] \quad (\text{linearity}) \end{aligned}$$

Estimators of regression parameters

Data $(X_1, Y_1), \dots, (X_N, Y_N)$ from $Y | X \sim \text{Normal}(\beta_0 + \beta_1 X, \sigma_\epsilon^2)$. How do we estimate β_1 ?

We know

$$\beta_1 = \frac{\text{cov}(X, Y)}{\sigma_X^2}.$$

To estimate $\text{cov}(X, Y)$ from data, it helps to write it in a different way. Start from

$$\begin{aligned} & E[(X - E[X])(Y - E[Y])] \\ &= E[XY - E[Y]X - E[X]Y + E[X]E[Y]] \quad (\text{expand}) \\ &= E[XY] - E[Y]E[X] - E[X]E[Y] + E[X]E[Y] \quad (\text{linearity}) \\ &= E[XY] - E[X]E[Y] = \text{cov}(X, Y). \end{aligned}$$

Estimators of regression parameters

Replace each expectation with a sample average, $E[\cdot] \rightarrow \frac{1}{N} \sum_i$ yields

$$\hat{\beta}_1 = \frac{\sum_i (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_i (X_i - \bar{X})^2}, \quad \hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}.$$

Estimators of regression parameters

Replace each expectation with a sample average, $E[\cdot] \rightarrow \frac{1}{N} \sum_i$ yields

$$\hat{\beta}_1 = \frac{\sum_i (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_i (X_i - \bar{X})^2}, \quad \hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}.$$

The residuals are $r_i = Y_i - (\hat{\beta}_0 + \hat{\beta}_1 X_i)$, and we define

$$\text{SSR} = \sum_i r_i^2 \quad (\text{sum of squared residuals}),$$

$$\text{SST} = \sum_i (Y_i - \bar{Y})^2 \quad (\text{total sum of squares}).$$

Estimators of regression parameters

Replace each expectation with a sample average, $E[\cdot] \rightarrow \frac{1}{N} \sum_i$ yields

$$\hat{\beta}_1 = \frac{\sum_i (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_i (X_i - \bar{X})^2}, \quad \hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}.$$

The residuals are $r_i = Y_i - (\hat{\beta}_0 + \hat{\beta}_1 X_i)$, and we define

$$\text{SSR} = \sum_i r_i^2 \quad (\text{sum of squared residuals}),$$

$$\text{SST} = \sum_i (Y_i - \bar{Y})^2 \quad (\text{total sum of squares}).$$

The estimators of the noise variance and of ρ^2 are

$$\hat{\sigma}_\epsilon^2 = \frac{\text{SSR}}{N-2} \approx \frac{\text{SSR}}{N} \quad \hat{\rho}^2 = R^2 = 1 - \frac{\text{SSR}}{\text{SST}}$$


($N-2$ is to make the estimator unbiased, but it makes a small difference once $N \approx 20$)

Least squares

$\hat{\beta}_0, \hat{\beta}_1$ give the line with the smallest SSR,

$$(\hat{\beta}_0, \hat{\beta}_1) = \underset{b_0, b_1}{\operatorname{argmin}} \sum_i (Y_i - b_0 - b_1 X_i)^2,$$

which is why they are called the least squares estimators. Other choices give other lines:




Same data, three lines (others in gray). LS and least abs. deviations use vertical residuals r_i ; PCA uses perpendicular distances d_i . Squaring makes LS sensitive to the outlier.


Two kinds of replicates

The sample distribution describes how $\hat{\beta}_0, \hat{\beta}_1$ vary over replicates of the data. For regression there are two ways to make a replicate:

conditional on X : same x 's, new y 's



not conditional on X : new x 's and new y 's




Each panel highlights one replicate; the other replicates in that row are in gray.


Two kinds of replicates

The sample distribution describes how $\hat{\beta}_0, \hat{\beta}_1$ vary over replicates of the data. For regression there are two ways to make a replicate:

conditional on X : same x 's, new y 's



not conditional on X : new x 's and new y 's



Each panel highlights one replicate; the other replicates in that row are in gray.

We focus on the sample distribution conditioned on $X_1, \dots, X_N$.

β_1 is an unbiased estimator

$$E[\hat{\beta}_1|X_1, \dots, X_N] = \frac{\sum_{i=1} (X_i - \bar{X}) E[(Y_i - \bar{Y})|X_1, \dots, X_N]}{\sum_{i=1} (X_i - \bar{X})^2}$$

and

$$E[(Y_i - \bar{Y})|X_1, \dots, X_N] = \beta_0 + \beta_1 X_i - \beta_0 - \beta_1 \bar{X} = \beta_1 (X_i - \bar{X})$$


so

$$E[\hat{\beta}_1|X_1, \dots, X_N] = \beta_1 \frac{\sum_{i=1} (X_i - \bar{X})^2}{\sum_{i=1} (X_i - \bar{X})^2} = \beta_1$$

Sample distribution of $\hat{\beta}_1$

For large N ,

$$\hat{\beta}_1 \sim \text{Normal}\left(\beta_1, \frac{\sigma_\epsilon^2}{N\sigma_X^2}\right)$$



$N = 20$, $\beta_1 = 2$, $\sigma_\epsilon = 1$. Top: one dataset, its fitted line (red), the 95% interval for the line (shaded) and the true line (dashed). Bottom: $\hat{\beta}_1$ over 5000 replicates with the same x 's, with the Normal approximation (dashed).

Back to the graders

- Recall in the graders example: $N = 68$ exams, $N_E = 23$ graded by E, and $\hat{\sigma}_\epsilon \approx 5$.
- Therefore $X \sim \text{Bernoulli}(p)$ with $p \approx N_E/N = 23/68 \approx 1/3$, and

$$\sigma_X^2 = p(1 - p) \approx 2/9 \implies N\sigma_X^2 \approx 68 \times \frac{2}{9} = \frac{136}{9} \approx 15.1.$$

Back to the graders

- Recall in the graders example: $N = 68$ exams, $N_E = 23$ graded by E, and $\hat{\sigma}_\epsilon \approx 5$.
- Therefore $X \sim \text{Bernoulli}(p)$ with $p \approx N_E/N = 23/68 \approx 1/3$, and

$$\sigma_X^2 = p(1-p) \approx 2/9 \implies N\sigma_X^2 \approx 68 \times \frac{2}{9} = \frac{136}{9} \approx 15.1.$$

- So $\hat{\beta}_1 \sim \text{Normal}(\beta_1, 1.65)$, since

$$\text{var}(\hat{\beta}_1) \approx \frac{\hat{\sigma}_\epsilon^2}{N\sigma_X^2} \approx \frac{25}{136/9} = \frac{225}{136} \approx 1.65, \quad \text{se}(\hat{\beta}_1) \approx \sqrt{1.65} \approx 1.29.$$

Back to the graders

- Recall in the graders example: $N = 68$ exams, $N_E = 23$ graded by E, and $\hat{\sigma}_\epsilon \approx 5$.
- Therefore $X \sim \text{Bernoulli}(p)$ with $p \approx N_E/N = 23/68 \approx 1/3$, and

$$\sigma_X^2 = p(1-p) \approx 2/9 \implies N\sigma_X^2 \approx 68 \times \frac{2}{9} = \frac{136}{9} \approx 15.1.$$

- So $\hat{\beta}_1 \sim \text{Normal}(\beta_1, 1.65)$, since

$$\text{var}(\hat{\beta}_1) \approx \frac{\hat{\sigma}_\epsilon^2}{N\sigma_X^2} \approx \frac{25}{136/9} = \frac{225}{136} \approx 1.65, \quad \text{se}(\hat{\beta}_1) \approx \sqrt{1.65} \approx 1.29.$$

- Compare with what we found before:
 - back of the envelope: $\text{var}(\hat{\beta}_1) \approx 3/2$, $\text{se}(\hat{\beta}_1) \approx 1.2$;
 - exact (Python): $\text{var}(\hat{\beta}_1) = 1.84$, $\text{se}(\hat{\beta}_1) \approx 1.36$.

All three agree up to rounding: the boxed formula gives the same answer as the difference of means.