

Lecture 11 (10/9)

Autoregressive models, regression to the mean, and the Hurwicz bias

Unit 3 so far:

- An estimator $\hat{\theta}$ is computed from data and varies from replicate to replicate.
- Its sample distribution ($\approx$ Normal) is used to obtain the 95% confidence interval $\hat{\theta} \pm 1.96 \text{se}(\hat{\theta})$ and the p-value for $H_0 : \theta = 0$.
- Bias-variance decomposition: $\text{MSE}_{\hat{\theta}} = \text{se}(\hat{\theta})^2 + \text{Bias}^2$.
- Least squares: $\hat{\beta}_0, \hat{\beta}_1$ minimize SSR; $\hat{\beta}_1 \approx \text{Normal}(\beta_1, \sigma_{\epsilon}^2 / (N\sigma_X^2))$.

Unit 3 so far:

- An estimator $\hat{\theta}$ is computed from data and varies from replicate to replicate.
- Its sample distribution ($\approx$ Normal) is used to obtain the 95% confidence interval $\hat{\theta} \pm 1.96 \text{se}(\hat{\theta})$ and the p-value for $H_0 : \theta = 0$.
- Bias-variance decomposition: $\text{MSE}_{\hat{\theta}} = \text{se}(\hat{\theta})^2 + \text{Bias}^2$.
- Least squares: $\hat{\beta}_0, \hat{\beta}_1$ minimize SSR; $\hat{\beta}_1 \approx \text{Normal}(\beta_1, \sigma_{\epsilon}^2 / (N\sigma_X^2))$.

Today:

- More practice working with (single-predictor) regression models
- Time series models.



p -values vs. R^2

Question: State whether each of the following scenarios is possible when we fit data to a linear regression model. The p -value (p_v) is for the null hypothesis $H_0 : \beta_1 = 0$ and assume you have $N > 3$ data points.

- (a) $p_v = 10^{-5}$ and $R^2 = 10^{-2}$.
- (b) $p_v = 0.5$ and $R^2 = 0.99$.
- (c) $\hat{\beta}_1 = 1000, \hat{\sigma}_\epsilon^2 = 10^{-5}$ and $R^2 = 10^{-4}$.



p -values vs. R^2 : answers

(a) and (c) yes.



p -values vs. R^2 : answers

(a) and (c) yes.



(b), no: for one predictor (see drill “The test statistic and R^2 ” in Section 3.6 of the notes.):

$$\hat{T}^2 = (N - 2) \frac{R^2}{1 - R^2}$$

so $R^2 = 0.99$ forces $\hat{T}^2 = 99(N - 2)$ and a tiny p_V .

Estimating parameters in Python

With `x`, `y` numpy arrays. Using the covariance:

```
C = np.cov(x, y)
beta1_hat = C[0, 1] / C[0, 0]          # cov(x, y) / var(x)
```

Here `C` is the 2×2 sample covariance matrix

$$C = \begin{pmatrix} \widehat{\text{var}}(X) & \widehat{\text{cov}}(X, Y) \\ \widehat{\text{cov}}(X, Y) & \widehat{\text{var}}(Y) \end{pmatrix}, \quad \widehat{\text{cov}}(X, Y) = \frac{1}{N-1} \sum_i (X_i - \bar{X})(Y_i - \bar{Y}),$$

so $C[0, 1] / C[0, 0] = \widehat{\text{cov}}(X, Y) / \widehat{\text{var}}(X) = \hat{\beta}_1$ (Python counts rows and columns from 0).

Estimating parameters in Python

We can also

```
beta1_hat = np.sum((x - np.mean(x)) * (y - np.mean(y))) \
              / np.sum((x - np.mean(x))**2)
```

Estimating parameters in Python

We can also







```
beta1_hat = np.sum((x - np.mean(x)) * (y - np.mean(y))) \
              / np.sum((x - np.mean(x))**2)
```

and the other estimators are:

```
beta0_hat = np.mean(y) - beta1_hat * np.mean(x)
r = y - (beta0_hat + beta1_hat * x)           # residuals
sigma2_hat = np.sum(r**2) / (len(y) - 2)
R2 = 1 - np.sum(r**2) / np.sum((y - np.mean(y))**2)
se_beta1 = np.sqrt(sigma2_hat / np.sum((x - np.mean(x))**2))
T = beta1_hat / se_beta1
p_v = 2 * (1 - norm.cdf(np.abs(T)))           # from scipy.stats
```







p -values vs. R^2

Guess R^2 and the p -value for $H_0 : \beta_1 = 0$ in each plot. No calculations needed!





p -values vs. R^2

Guess R^2 and the p -value for $H_0 : \beta_1 = 0$ in each plot. No calculations needed!





Time series data

Here $Y_t = \log(\text{price on day } t)$.



Time series data



Here $Y_t = \log(\text{price on day } t)$.



Goal: Predict the price tomorrow, and the next day ...

Time series data

Here $Y_t = \log(\text{price on day } t)$.




Goal: Predict the price tomorrow, and the next day ...

- Could try treating time as predictors, but this doesn't really make sense: Each time doesn't independently produce price. In other words, the data points (t, Y_t) and $(t + 1, Y_{t+1})$ are not independent.

Time series data

Here $Y_t = \log(\text{price on day } t)$.



Goal: Predict the price tomorrow, and the next day ...

- Could try treating time as predictors, but this doesn't really make sense: Each time doesn't independently produce price. In other words, the data points (t, Y_t) and $(t + 1, Y_{t+1})$ are not independent.
- Instead, we want to capture the fact that prices depend on the previous prices.

Autoregressive models

An autoregressive model (AR(1)) is a linear regression model where the predictor is the previous value of the response:

$$Y_t \mid Y_{t-1} \sim \text{Normal}(\beta_0 + \beta_1 Y_{t-1}, \sigma_\epsilon^2), \quad \text{i.e.} \quad Y_t = \beta_0 + \beta_1 Y_{t-1} + \epsilon_t.$$

Autoregressive models

An autoregressive model (AR(1)) is a linear regression model where the predictor is the previous value of the response:

$$Y_t \mid Y_{t-1} \sim \text{Normal}(\beta_0 + \beta_1 Y_{t-1}, \sigma_\epsilon^2), \quad \text{i.e.} \quad Y_t = \beta_0 + \beta_1 Y_{t-1} + \epsilon_t.$$

The data for the regression are the pairs

$$(X_1, Y_1), (X_2, Y_2), \dots, (X_N, Y_N) = (Y_1, Y_2), (Y_2, Y_3), \dots, (Y_{T-1}, Y_T),$$


Autoregressive models

An autoregressive model (AR(1)) is a linear regression model where the predictor is the previous value of the response:

$$Y_t \mid Y_{t-1} \sim \text{Normal}(\beta_0 + \beta_1 Y_{t-1}, \sigma_\epsilon^2), \quad \text{i.e.} \quad Y_t = \beta_0 + \beta_1 Y_{t-1} + \epsilon_t.$$

The data for the regression are the pairs

$$(X_1, Y_1), (X_2, Y_2), \dots, (X_N, Y_N) = (Y_1, Y_2), (Y_2, Y_3), \dots, (Y_{T-1}, Y_T),$$



Stationary process

- A series is stationary if Y_t and Y_{t-1} have the same mean μ_Y and the same standard deviation σ_Y . Then

$$\beta_1 = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y^2} = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y \sigma_Y} = \rho,$$

Stationary process

- A series is stationary if Y_t and Y_{t-1} have the same mean μ_Y and the same standard deviation σ_Y . Then

$$\beta_1 = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y^2} = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y \sigma_Y} = \rho,$$

- Taking expectations,

$$\mu = \beta_0 + \rho\mu \implies \beta_0 = (1 - \rho)\mu$$

- Therefore

$$Y_t = \mu(1 - \rho) + \rho Y_t + \epsilon_t.$$

Stationary process

- A series is stationary if Y_t and Y_{t-1} have the same mean μ_Y and the same standard deviation σ_Y . Then

$$\beta_1 = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y^2} = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y \sigma_Y} = \rho,$$

- Taking expectations,

$$\mu = \beta_0 + \rho\mu \implies \beta_0 = (1 - \rho)\mu$$

- Therefore


$$Y_t = \mu(1 - \rho) + \rho Y_t + \epsilon_t.$$

- Taking the variance

$$\sigma_Y^2 = \sigma_Y^2 \rho^2 + \sigma_\epsilon^2 \implies \sigma_Y^2 = \frac{\sigma_\epsilon^2}{1 - \rho^2}$$

Simulated AR(1) models

$Y_{t+1} = \rho Y_t + \epsilon_{t+1}$ with $\epsilon \sim \text{Normal}(0, 1)$. Top: Y_{t+1} vs. Y_t for one series;
bottom: a few series of 100 steps.



Estimating ρ : The Hurwicz bias

Take $\mu = 0$ for simplicity. The mean of the least squares estimator is

$$E[\hat{\rho}] = E \left[\frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2} \right]$$

Estimating ρ : The Hurwicz bias

Take $\mu = 0$ for simplicity. The mean of the least squares estimator is

$$E[\hat{\rho}] = E \left[\frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2} \right] = E \left[\frac{\sum_t Y_{t-1} (\rho Y_{t-1} + \epsilon_t)}{\sum_t Y_{t-1}^2} \right]$$

Estimating ρ : The Hurwicz bias

Take $\mu = 0$ for simplicity. The mean of the least squares estimator is

$$\begin{aligned} E[\hat{\rho}] &= E \left[\frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2} \right] = E \left[\frac{\sum_t Y_{t-1} (\rho Y_{t-1} + \epsilon_t)}{\sum_t Y_{t-1}^2} \right] \\ &= \rho + E \left[\frac{\sum_t Y_{t-1} \epsilon_t}{\sum_t Y_{t-1}^2} \right] \end{aligned}$$

Estimating ρ : The Hurwicz bias

Take $\mu = 0$ for simplicity. The mean of the least squares estimator is

$$\begin{aligned} E[\hat{\rho}] &= E \left[\frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2} \right] = E \left[\frac{\sum_t Y_{t-1} (\rho Y_{t-1} + \epsilon_t)}{\sum_t Y_{t-1}^2} \right] \\ &= \rho + E \left[\frac{\sum_t Y_{t-1} \epsilon_t}{\sum_t Y_{t-1}^2} \right] \underset{\text{Hurwicz (1950)}}{\approx} \rho - \frac{2\rho}{N}. \end{aligned}$$

Estimating ρ : The Hurwicz bias

Take $\mu = 0$ for simplicity. The mean of the least squares estimator is

$$\begin{aligned} E[\hat{\rho}] &= E \left[\frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2} \right] = E \left[\frac{\sum_t Y_{t-1} (\rho Y_{t-1} + \epsilon_t)}{\sum_t Y_{t-1}^2} \right] \\ &= \rho + E \left[\frac{\sum_t Y_{t-1} \epsilon_t}{\sum_t Y_{t-1}^2} \right] \underset{\text{Hurwicz (1950)}}{\approx} \rho - \frac{2\rho}{N}. \end{aligned}$$


In ordinary regression, the noise ϵ_i is independent of *all* of the predictors, so conditional on X the second term has mean 0 and $\hat{\beta}_1$ is unbiased.

Estimating ρ : The Hurwicz bias

Take $\mu = 0$ for simplicity. The mean of the least squares estimator is

$$\begin{aligned} E[\hat{\rho}] &= E \left[\frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2} \right] = E \left[\frac{\sum_t Y_{t-1} (\rho Y_{t-1} + \epsilon_t)}{\sum_t Y_{t-1}^2} \right] \\ &= \rho + E \left[\frac{\sum_t Y_{t-1} \epsilon_t}{\sum_t Y_{t-1}^2} \right] \underset{\text{Hurwicz (1950)}}{\approx} \rho - \frac{2\rho}{N}. \end{aligned}$$

In ordinary regression, the noise ϵ_i is independent of *all* of the predictors, so conditional on X the second term has mean 0 and $\hat{\beta}_1$ is unbiased.






When $\rho \approx 1$: The random walk

If $\rho \approx 1$, then $\sigma_Y \approx \infty$. How can this be?

$$Y_{t+1} = Y_t + \epsilon_t = Y_{t-1} + \epsilon_{t-1} + \epsilon_t = \cdots = \sum_{s \leq t} \epsilon_s$$

And hence the marginal distribution of Y_t is

$$Y_{t+1} \sim \text{Normal}(0, \sigma_\epsilon^2 t)$$



How fast do correlations decay?

The autocorrelation at lag k is $\rho_k = \text{corr}(Y_t, Y_{t+k})$. For a stationary AR(1) series with mean 0, $Y_{t+1} = \rho Y_t + \epsilon_{t+1}$:

How fast do correlations decay?

The autocorrelation at lag k is $\rho_k = \text{corr}(Y_t, Y_{t+k})$. For a stationary AR(1) series with mean 0, $Y_{t+1} = \rho Y_t + \epsilon_{t+1}$:

$$Y_{t+1} = \rho Y_t + \epsilon_{t+1}$$

How fast do correlations decay?

The autocorrelation at lag k is $\rho_k = \text{corr}(Y_t, Y_{t+k})$. For a stationary AR(1) series with mean 0, $Y_{t+1} = \rho Y_t + \epsilon_{t+1}$:

$$Y_{t+1} = \rho Y_t + \epsilon_{t+1}$$

$$Y_{t+2} = \rho Y_{t+1} + \epsilon_{t+2} = \rho^2 Y_t + \rho \epsilon_{t+1} + \epsilon_{t+2}$$

How fast do correlations decay?

The autocorrelation at lag k is $\rho_k = \text{corr}(Y_t, Y_{t+k})$. For a stationary AR(1) series with mean 0, $Y_{t+1} = \rho Y_t + \epsilon_{t+1}$:

$$Y_{t+1} = \rho Y_t + \epsilon_{t+1}$$

$$Y_{t+2} = \rho Y_{t+1} + \epsilon_{t+2} = \rho^2 Y_t + \rho \epsilon_{t+1} + \epsilon_{t+2}$$

$$\vdots$$

$$Y_{t+k} = \rho^k Y_t + \underbrace{\rho^{k-1} \epsilon_{t+1} + \cdots + \rho \epsilon_{t+k-1} + \epsilon_{t+k}}_{\text{future noise, independent of } Y_t}$$

How fast do correlations decay?

The autocorrelation at lag k is $\rho_k = \text{corr}(Y_t, Y_{t+k})$. For a stationary AR(1) series with mean 0, $Y_{t+1} = \rho Y_t + \epsilon_{t+1}$:

$$Y_{t+1} = \rho Y_t + \epsilon_{t+1}$$

$$Y_{t+2} = \rho Y_{t+1} + \epsilon_{t+2} = \rho^2 Y_t + \rho \epsilon_{t+1} + \epsilon_{t+2}$$

$$\vdots$$

$$Y_{t+k} = \rho^k Y_t + \underbrace{\rho^{k-1} \epsilon_{t+1} + \cdots + \rho \epsilon_{t+k-1} + \epsilon_{t+k}}_{\text{future noise, independent of } Y_t}$$

So $\text{cov}(Y_t, Y_{t+k}) = \rho^k \text{var}(Y_t) + 0 = \rho^k \sigma_Y^2$, and since $\text{var}(Y_{t+k}) = \text{var}(Y_t) = \sigma_Y^2$,

$$\boxed{\rho_k = \frac{\text{cov}(Y_t, Y_{t+k})}{\sigma_Y \sigma_Y} = \rho^k}$$

How fast do correlations decay?

The autocorrelation at lag k is $\rho_k = \text{corr}(Y_t, Y_{t+k})$. For a stationary AR(1) series with mean 0, $Y_{t+1} = \rho Y_t + \epsilon_{t+1}$:

$$Y_{t+1} = \rho Y_t + \epsilon_{t+1}$$

$$Y_{t+2} = \rho Y_{t+1} + \epsilon_{t+2} = \rho^2 Y_t + \rho \epsilon_{t+1} + \epsilon_{t+2}$$

$$\vdots$$

$$Y_{t+k} = \rho^k Y_t + \underbrace{\rho^{k-1} \epsilon_{t+1} + \cdots + \rho \epsilon_{t+k-1} + \epsilon_{t+k}}_{\text{future noise, independent of } Y_t}$$


So $\text{cov}(Y_t, Y_{t+k}) = \rho^k \text{var}(Y_t) + 0 = \rho^k \sigma_Y^2$, and since $\text{var}(Y_{t+k}) = \text{var}(Y_t) = \sigma_Y^2$,

$$\boxed{\rho_k = \frac{\text{cov}(Y_t, Y_{t+k})}{\sigma_Y \sigma_Y} = \rho^k}$$

The correlation decays *exponentially* in the lag: after $k = 10$ steps, $0.5^{10} \approx 0.001$, $0.9^{10} \approx 0.35$, $0.99^{10} \approx 0.90$.


Autocorrelation functions

Sample version:
$$\hat{\rho}_k = \frac{\sum_t (Y_t - \bar{Y})(Y_{t+k} - \bar{Y})}{\sum_t (Y_t - \bar{Y})^2}.$$



Autocorrelation functions

Sample version:
$$\hat{\rho}_k = \frac{\sum_t (Y_t - \bar{Y})(Y_{t+k} - \bar{Y})}{\sum_t (Y_t - \bar{Y})^2}.$$




- For the simulated series, $\hat{\rho}_k$ follows ρ^k .
- The log price itself stays correlated for a long time; its daily changes $Y_t - Y_{t-1}$ are almost uncorrelated.

RW vs. linear extrapolation



Fit two models to $N = 180$ days of log price and forecast the next 90 days:

- **linear in time**: Fitting $Y_t = \beta_0 + \beta_1 t + \epsilon_t$ and extrapolating. Shaded area is the 95% prediction band, including both ϵ noise and uncertainty in β_1 .
- **random walk**: Forecast $Y_{t+h} \approx Y_t$, with 95% band $\pm 1.96 \hat{\sigma}_\epsilon \sqrt{h}$.




Prices vs. log prices

So far $Y_t = \log P_t$, where P_t is the price. On a linear scale:



Prices vs. log prices

So far $Y_t = \log P_t$, where P_t is the price. On a linear scale:




If the log price is (close to) a random walk, $Y_{t+1} = Y_t + \epsilon_{t+1}$, then

$$P_{t+1} = e^{Y_{t+1}} = P_t e^{\epsilon_{t+1}} \approx P_t (1 + \epsilon_{t+1}) :$$

each day the price is *multiplied* by a random factor (a multiplicative process).

Prices vs. log prices

So far $Y_t = \log P_t$, where P_t is the price. On a linear scale:



If the log price is (close to) a random walk, $Y_{t+1} = Y_t + \epsilon_{t+1}$, then

$$P_{t+1} = e^{Y_{t+1}} = P_t e^{\epsilon_{t+1}} \approx P_t (1 + \epsilon_{t+1}) :$$

each day the price is *multiplied* by a random factor (a multiplicative process). So $P_{t+1} - P_t \approx P_t \epsilon_{t+1}$: the daily change is proportional to the current price. The relative volatility is constant ($\hat{\sigma}_\epsilon \approx 0.04$, about 4% a day), but in dollars the swings grow with the price: a typical daily move was about \$3 in 2015 and about \$900 in the last year.