

Lecture 11 (10/9)

Autoregressive models, regression to the mean, and the Hurwicz bias

- More practice working with (single-predictor) regression models
- Time series models (the fun part)

p -values vs. R^2







Question:

Is each of the following possible? If so, how? (The p -value is for the null hypothesis $H_0 : \beta_1 = 0$.)

- A very small p -value, but $R^2 \approx 0$.
- A p -value $\approx 1/2$, but R^2 close to 1.
- A massive β_1 but tiny R^2 .

p -values vs. R^2


Guess R^2 and the p -value for $H_0 : \beta_1 = 0$ in each plot. No calculations needed!



Time series data

The price of bitcoin each day since December 2014, on a log scale:


$$Y_t = \log_{10}(\text{price on day } t).$$



Time series data

The price of bitcoin each day since December 2014, on a log scale:

$$Y_t = \log_{10}(\text{price on day } t).$$



How do we model this?

Conditioning on the history

In general, the next value depends on the whole history:

$$Y_t \mid Y_{t-1}, Y_{t-2}, \dots, Y_1.$$

Conditioning on the history

In general, the next value depends on the whole history:

$$Y_t \mid Y_{t-1}, Y_{t-2}, \dots, Y_1.$$

The simplest assumption: the next day depends only on the current day,

$$Y_t \mid Y_{t-1}, Y_{t-2}, \dots, Y_1 \text{ has the same distribution as } Y_t \mid Y_{t-1}.$$


Conditioning on the history

In general, the next value depends on the whole history:

$$Y_t \mid Y_{t-1}, Y_{t-2}, \dots, Y_1.$$

The simplest assumption: the next day depends only on the current day,

$Y_t \mid Y_{t-1}, Y_{t-2}, \dots, Y_1$ has the same distribution as $Y_t \mid Y_{t-1}$.



Autoregressive models

An autoregressive model (AR(1)) is a linear regression model where the predictor is the previous value of the response:

$$Y_t \mid Y_{t-1} \sim \text{Normal}(\beta_0 + \beta_1 Y_{t-1}, \sigma_\epsilon^2), \quad \text{i.e.} \quad Y_t = \beta_0 + \beta_1 Y_{t-1} + \epsilon_t.$$

Autoregressive models

An autoregressive model (AR(1)) is a linear regression model where the predictor is the previous value of the response:

$$Y_t \mid Y_{t-1} \sim \text{Normal}(\beta_0 + \beta_1 Y_{t-1}, \sigma_\epsilon^2), \quad \text{i.e.} \quad Y_t = \beta_0 + \beta_1 Y_{t-1} + \epsilon_t.$$

The data for the regression are the pairs

$$(X, Y) = (Y_1, Y_2), (Y_2, Y_3), \dots, (Y_{T-1}, Y_T),$$

so we can estimate β_0, β_1 with least squares exactly as before.

Autoregressive models

An autoregressive model (AR(1)) is a linear regression model where the predictor is the previous value of the response:

$$Y_t \mid Y_{t-1} \sim \text{Normal}(\beta_0 + \beta_1 Y_{t-1}, \sigma_\epsilon^2), \quad \text{i.e.} \quad Y_t = \beta_0 + \beta_1 Y_{t-1} + \epsilon_t.$$

The data for the regression are the pairs

$$(X, Y) = (Y_1, Y_2), (Y_2, Y_3), \dots, (Y_{T-1}, Y_T),$$

so we can estimate β_0, β_1 with least squares exactly as before.

Example

For bitcoin, least squares gives $\hat{\beta}_1 \approx 1.000$. What does this say about tomorrow's value?

The coefficient is a correlation

Y_t and Y_{t-1} are the *same quantity*, with the same units, so β_1 must be dimensionless.

The coefficient is a correlation

Y_t and Y_{t-1} are the *same quantity*, with the same units, so β_1 must be dimensionless.

Suppose the series is stationary: Y_t and Y_{t-1} have the same mean μ and the same standard deviation σ_Y . Then

$$\beta_1 = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y^2} = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y \sigma_Y} = \rho,$$

the correlation between consecutive values.

The coefficient is a correlation

Y_t and Y_{t-1} are the *same quantity*, with the same units, so β_1 must be dimensionless.

Suppose the series is stationary: Y_t and Y_{t-1} have the same mean μ and the same standard deviation σ_Y . Then

$$\beta_1 = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y^2} = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y \sigma_Y} = \rho,$$

the correlation between consecutive values. Taking expectations, $\mu = \beta_0 + \rho\mu$, so $\beta_0 = (1 - \rho)\mu$ and

$$Y_t - \mu = \rho(Y_{t-1} - \mu) + \epsilon_t.$$

The coefficient is a correlation

Y_t and Y_{t-1} are the *same quantity*, with the same units, so β_1 must be dimensionless.

Suppose the series is stationary: Y_t and Y_{t-1} have the same mean μ and the same standard deviation σ_Y . Then

$$\beta_1 = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y^2} = \frac{\text{cov}(Y_{t-1}, Y_t)}{\sigma_Y \sigma_Y} = \rho,$$

the correlation between consecutive values. Taking expectations, $\mu = \beta_0 + \rho\mu$, so $\beta_0 = (1 - \rho)\mu$ and

$$Y_t - \mu = \rho(Y_{t-1} - \mu) + \epsilon_t.$$

If $|\rho| < 1$, each step shrinks the deviation from the mean by a factor ρ :
regression to the mean.

Regression to the mean: some history

- Francis Galton, *Regression towards mediocrity in hereditary stature* (1886): compared the heights of about 900 adult children with those of their parents.

Regression to the mean: some history

- Francis Galton, *Regression towards mediocrity in hereditary stature* (1886): compared the heights of about 900 adult children with those of their parents.
- Children of unusually tall parents were tall, but on average *less* tall than their parents; children of short parents were less short. Galton called this “regression,” and that is where the name of this whole subject comes from.

Regression to the mean: some history


- Francis Galton, *Regression towards mediocrity in hereditary stature* (1886): compared the heights of about 900 adult children with those of their parents.
- Children of unusually tall parents were tall, but on average *less* tall than their parents; children of short parents were less short. Galton called this “regression,” and that is where the name of this whole subject comes from.
- In standardized units, $E[\text{child} \mid \text{parent}] = \rho \cdot \text{parent}$ with $\rho < 1$. Nothing biological is pulling heights back; it is what happens whenever the correlation is less than 1.

Regression to the mean: some history

- Francis Galton, *Regression towards mediocrity in hereditary stature* (1886): compared the heights of about 900 adult children with those of their parents.
- Children of unusually tall parents were tall, but on average *less* tall than their parents; children of short parents were less short. Galton called this “regression,” and that is where the name of this whole subject comes from.
- In standardized units, $E[\text{child} \mid \text{parent}] = \rho \cdot \text{parent}$ with $\rho < 1$. Nothing biological is pulling heights back; it is what happens whenever the correlation is less than 1.
- It shows up everywhere: the “sophomore slump” after a great rookie season, or flight instructors who conclude that praise makes pilots worse and criticism makes them better (Kahneman).


Simulations with different ρ

$Y_t = \rho Y_{t-1} + \epsilon_t$ with $\epsilon_t \sim \text{Normal}(0, 1)$, the same noise in every panel:



Simulations with different ρ

$Y_t = \rho Y_{t-1} + \epsilon_t$ with $\epsilon_t \sim \text{Normal}(0, 1)$, the same noise in every panel:



Larger ρ : longer excursions away from the mean. Negative ρ : the series flips sign from step to step.

The Hurwicz bias

Take $\mu = 0$ and fit $Y_t = \rho Y_{t-1} + \epsilon_t$ by least squares (no intercept):

$$\hat{\rho} = \frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2}$$

The Hurwicz bias

Take $\mu = 0$ and fit $Y_t = \rho Y_{t-1} + \epsilon_t$ by least squares (no intercept):

$$\hat{\rho} = \frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2} = \frac{\sum_t Y_{t-1} (\rho Y_{t-1} + \epsilon_t)}{\sum_t Y_{t-1}^2}$$

The Hurwicz bias

Take $\mu = 0$ and fit $Y_t = \rho Y_{t-1} + \epsilon_t$ by least squares (no intercept):

$$\hat{\rho} = \frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2} = \frac{\sum_t Y_{t-1} (\rho Y_{t-1} + \epsilon_t)}{\sum_t Y_{t-1}^2} = \rho + \frac{\sum_t Y_{t-1} \epsilon_t}{\sum_t Y_{t-1}^2}.$$

The Hurwicz bias

Take $\mu = 0$ and fit $Y_t = \rho Y_{t-1} + \epsilon_t$ by least squares (no intercept):

$$\hat{\rho} = \frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2} = \frac{\sum_t Y_{t-1} (\rho Y_{t-1} + \epsilon_t)}{\sum_t Y_{t-1}^2} = \rho + \frac{\sum_t Y_{t-1} \epsilon_t}{\sum_t Y_{t-1}^2}.$$

In ordinary regression, the noise ϵ_i is independent of *all* of the predictors, so conditional on X the second term has mean 0 and $\hat{\beta}_1$ is unbiased.

The Hurwicz bias

Take $\mu = 0$ and fit $Y_t = \rho Y_{t-1} + \epsilon_t$ by least squares (no intercept):

$$\hat{\rho} = \frac{\sum_t Y_{t-1} Y_t}{\sum_t Y_{t-1}^2} = \frac{\sum_t Y_{t-1} (\rho Y_{t-1} + \epsilon_t)}{\sum_t Y_{t-1}^2} = \rho + \frac{\sum_t Y_{t-1} \epsilon_t}{\sum_t Y_{t-1}^2}.$$

In ordinary regression, the noise ϵ_i is independent of *all* of the predictors, so conditional on X the second term has mean 0 and $\hat{\beta}_1$ is unbiased.

Here the predictors are $Y_1, \dots, Y_{T-1}$ themselves. ϵ_t is independent of the past (Y_{t-1}), but it is part of every future value:

$$Y_t = \rho Y_{t-1} + \epsilon_t, \quad Y_{t+1} = \rho Y_t + \epsilon_{t+1}, \dots$$

So the predictors are correlated with the noise, and we can no longer argue the second term averages to 0.

The Hurwicz bias

Why is it negative? Each term $Y_{t-1}\epsilon_t$ in the numerator has mean 0, but the numerator and denominator move together:

- when ϵ_t has the same sign as Y_{t-1} , the numerator goes up *and* $|Y_t|$ is large, which makes the denominator $\sum_t Y_{t-1}^2$ large;
- so positive contributions get divided by larger numbers than negative ones.

The Hurwicz bias

Why is it negative? Each term $Y_{t-1}\epsilon_t$ in the numerator has mean 0, but the numerator and denominator move together:

- when ϵ_t has the same sign as Y_{t-1} , the numerator goes up *and* $|Y_t|$ is large, which makes the denominator $\sum_t Y_{t-1}^2$ large;
- so positive contributions get divided by larger numbers than negative ones.

A longer calculation (Hurwicz, 1950) gives, for large T ,

$$E[\hat{\rho}] - \rho \approx -\frac{2\rho}{T}.$$


The Hurwicz bias

Why is it negative? Each term $Y_{t-1}\epsilon_t$ in the numerator has mean 0, but the numerator and denominator move together:

- when ϵ_t has the same sign as Y_{t-1} , the numerator goes up *and* $|Y_t|$ is large, which makes the denominator $\sum_t Y_{t-1}^2$ large;
- so positive contributions get divided by larger numbers than negative ones.

A longer calculation (Hurwicz, 1950) gives, for large T ,

$$E[\hat{\rho}] - \rho \approx -\frac{2\rho}{T}.$$



Biased toward 0 at every T , but consistent.

How fast do correlations decay?

The autocorrelation at lag k is $\rho_k = \text{corr}(Y_t, Y_{t+k})$. For AR(1), iterating $Y_{t+1} = \rho Y_t + \epsilon_{t+1}$ gives

$$Y_{t+k} = \rho^k Y_t + (\text{noise independent of } Y_t) \quad \implies \quad \rho_k = \rho^k.$$

How fast do correlations decay?

The autocorrelation at lag k is $\rho_k = \text{corr}(Y_t, Y_{t+k})$. For AR(1), iterating $Y_{t+1} = \rho Y_t + \epsilon_{t+1}$ gives

$$Y_{t+k} = \rho^k Y_t + (\text{noise independent of } Y_t) \quad \implies \quad \rho_k = \rho^k.$$


Sample version: $\hat{\rho}_k = \frac{\sum_t (Y_t - \bar{Y})(Y_{t+k} - \bar{Y})}{\sum_t (Y_t - \bar{Y})^2}.$

How fast do correlations decay?

The autocorrelation at lag k is $\rho_k = \text{corr}(Y_t, Y_{t+k})$. For AR(1), iterating $Y_{t+1} = \rho Y_t + \epsilon_{t+1}$ gives

$$Y_{t+k} = \rho^k Y_t + (\text{noise independent of } Y_t) \quad \implies \quad \rho_k = \rho^k.$$





Sample version:
$$\hat{\rho}_k = \frac{\sum_t (Y_t - \bar{Y})(Y_{t+k} - \bar{Y})}{\sum_t (Y_t - \bar{Y})^2}.$$



The log price itself stays correlated for a long time; its daily changes $Y_t - Y_{t-1}$ are almost uncorrelated.


Climate data

Monthly CO₂ concentration at Mauna Loa, Hawaii (NOAA):



Climate data

Monthly CO₂ concentration at Mauna Loa, Hawaii (NOAA):



- Is an AR(1) model reasonable here? What does $\hat{\rho}_k \approx 1$ for every k tell you?
- What does the autocorrelation of the month-to-month changes show?
- In Assignment 2 you will make the same kind of plot for daily casualty counts. A series with a strong trend (like a cumulative count) is strongly correlated with time and with itself, whatever the day-to-day details.